

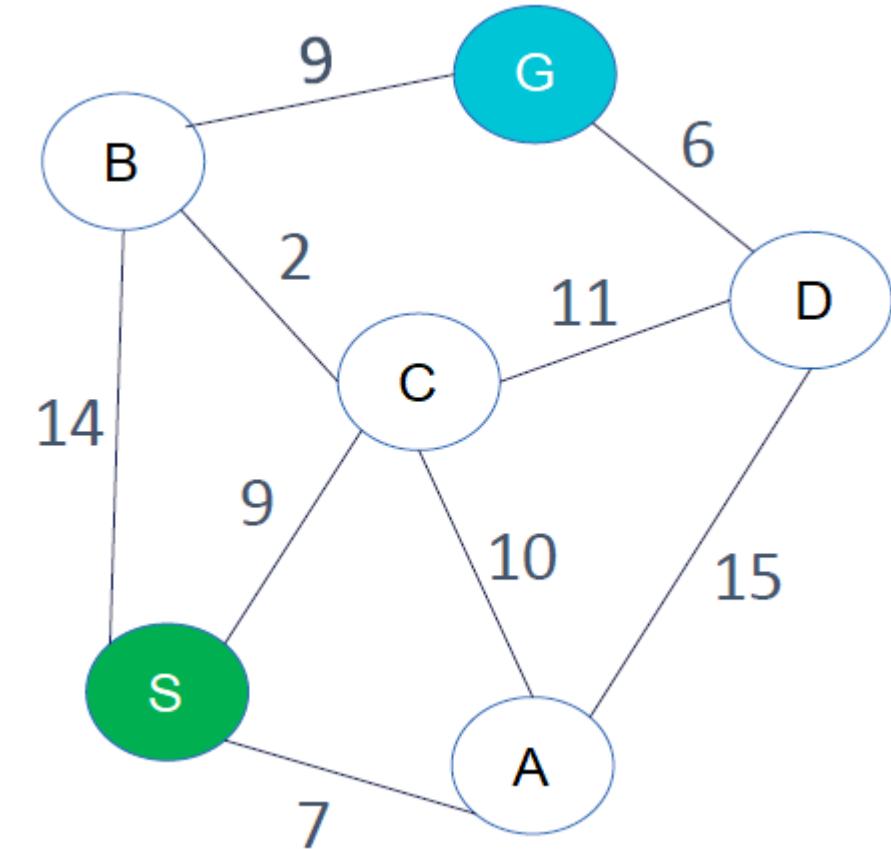
Introduction to Artificial Intelligence in Games

Lecture 4

Uniform-Cost Search

Task

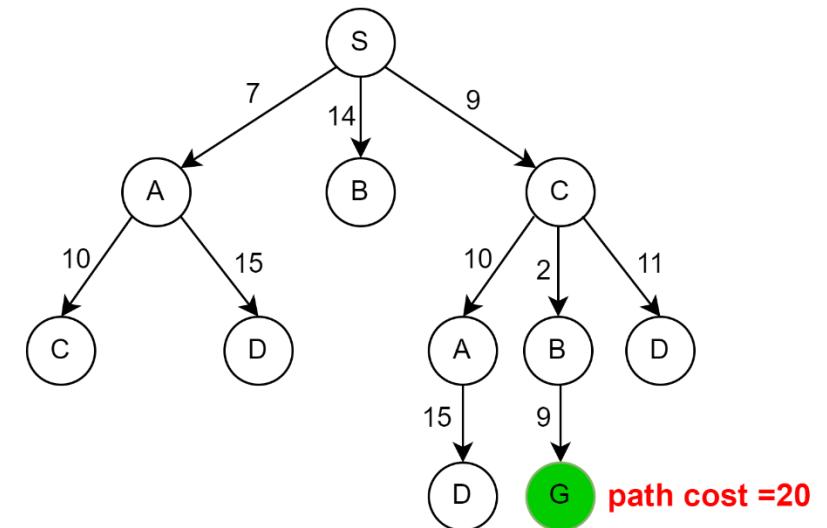
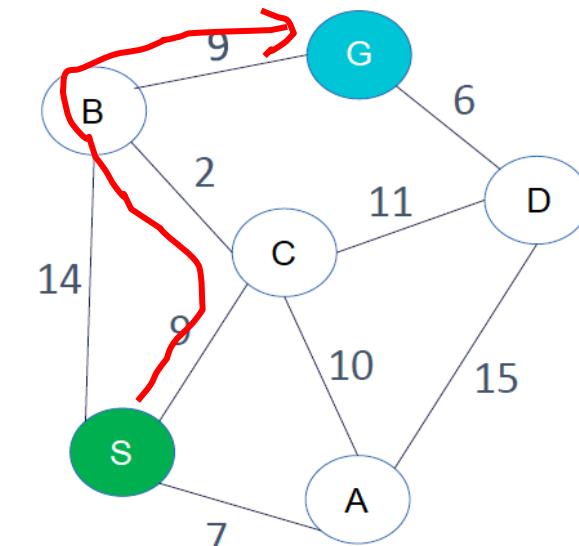
- Construct the search tree and find the optimum path from S to G using the uniform-cost search algorithm.



Solution

Expanded nodes list (sort then expand)

- S0 {A7, C9, B14 }
- A7 { C9, B14, C17, D22 }
- C9 { B11, B14, A19, D20, D22 }
- B11 { A19, G20, D20 }
- A19 { G20, D20, D34 }
- G20
- Solution path found is S → C → B → G, cost= 20



Informed Search Methods

Informed Search Algorithms

- In informed search algorithm, we have some knowledge about the goal, such as how far we are from the goal (straight line distance), how many wrong-placed tiles in 8-puzzle game, etc.
- This knowledge help agents to explore less to the search space and find more efficiently the goal node.
- Informed search algorithm uses the idea of heuristic, so it is also called **Heuristic search**.
- In a heuristic search, each state is assigned a “heuristic value” (h -value) that the search uses in selecting the “best” next step.
- Heuristic $h(n)$ gives an estimate of the distance from n to the goal.
- $h(n)=0$ for goal nodes

Greedy Best-first Search

- The basic idea of best first search is similar to uniform cost search.
- The only difference is that the “cost” (or in general, the evaluation function) is estimated based on some heuristics $h(n)$ which represents how far we are from the goal, rather than the accumulated cost $g(n)$ calculated during search.
- In the best first search algorithm, we expand the node which is closest to the goal node.
- Similar to the uniform cost search, the greedy best-first algorithm is implemented by the priority queue, with the node that has the lowest heuristic value h is processed first.

Greedy Best-first Search

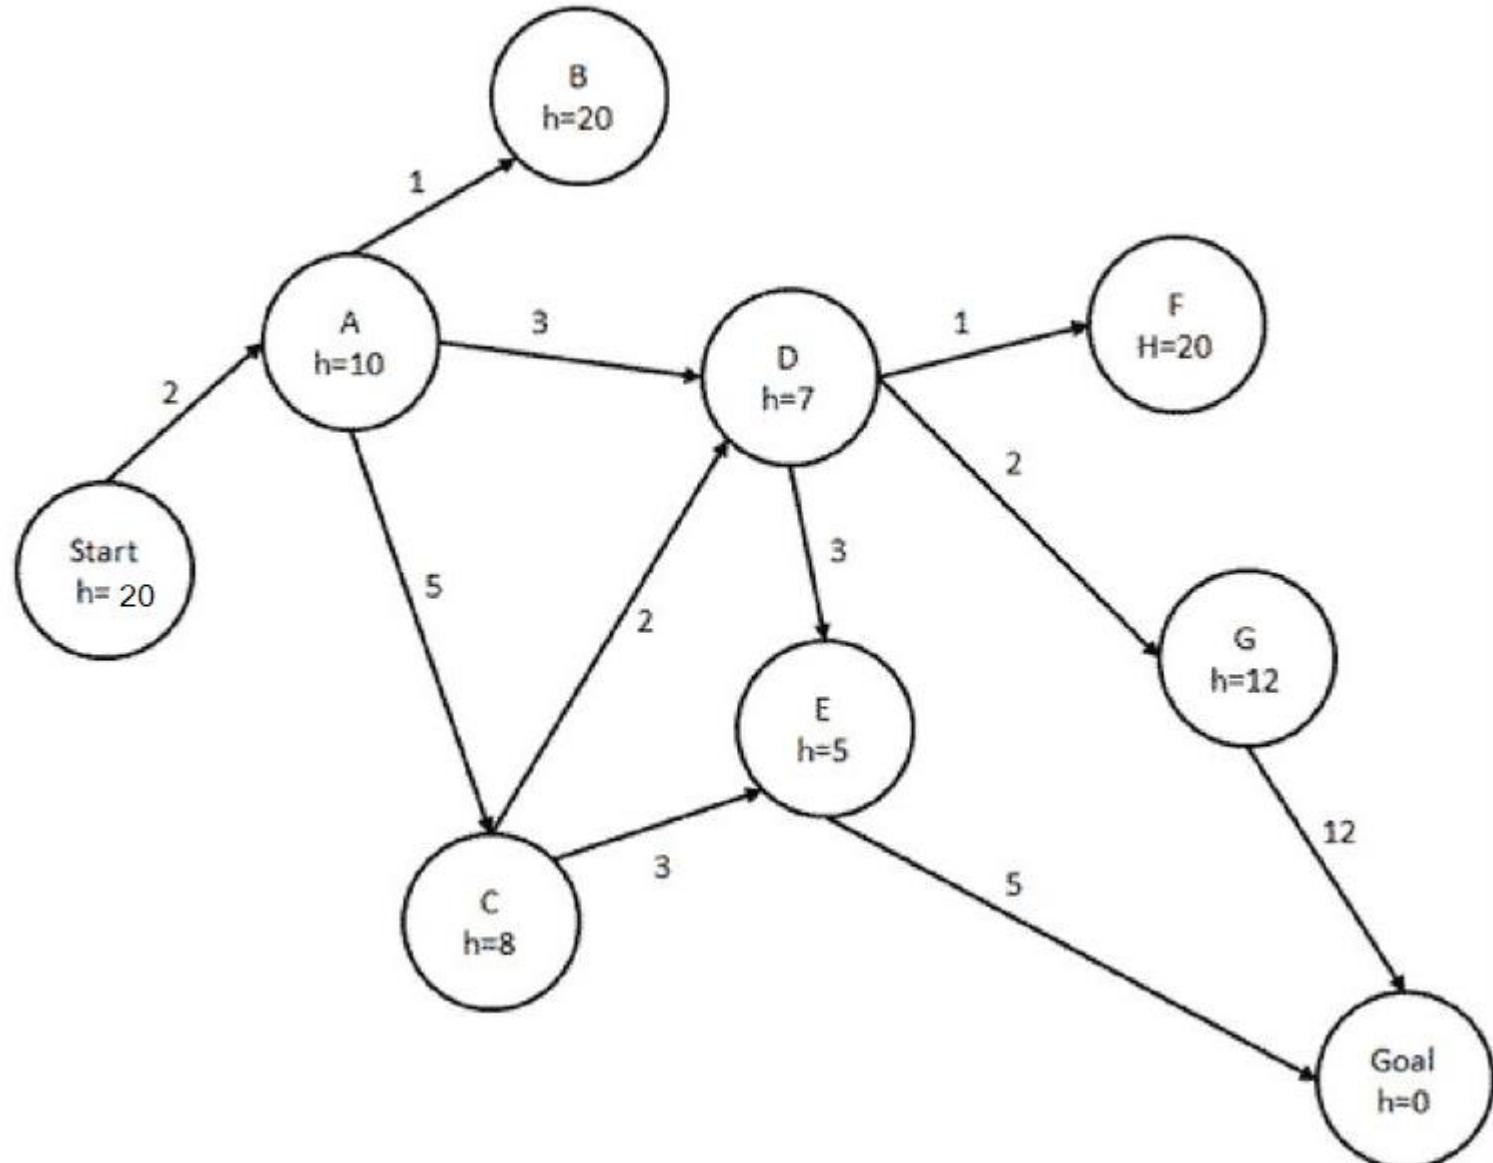
Steps

- Enqueue nodes by their heuristic value.
- That is, let $h(n) = \text{cost (straight line distance)} \text{ of the path from the current node } n \text{ to the goal node.}$
- Sort nodes by ascending value of h
- If there is more than one path to the same node, keep only the path with the lowest cost in the queue.
- Stop when the goal is the first node in the queue
Else, expand nodes starting from the first node in the queue
- Ignore back paths to the start node.
- Ignore back paths to the parent node.

Greedy Best-first Search

Example

- Find the path from the start node to the goal node using the greedy best-first search.
- The heuristic value $h(n)$ for each node is shown on the node.

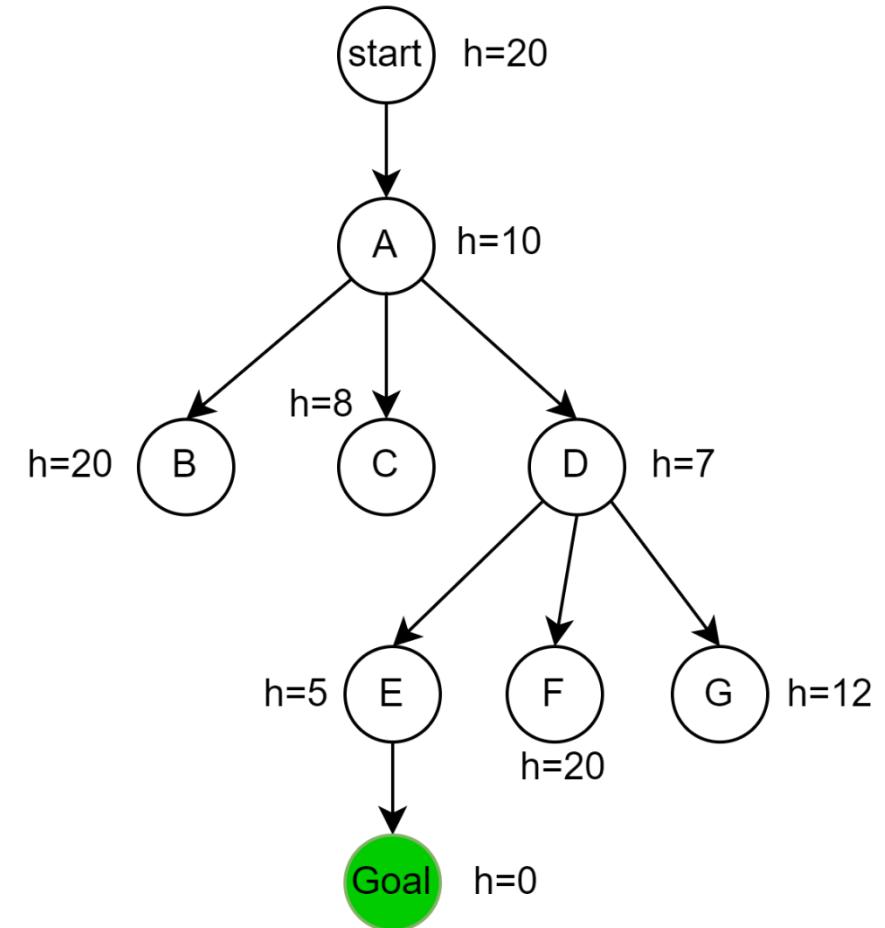


Greedy Best-first Search

Solution

Expanded nodes list (sort then expand)

- Start { A10 }
- A10 { D7, C8, B20}
- D7 { E5, C8, G12, B20, F20 }
- E5 { Goal }
- Goal
- Solution path found is Start → A → D → E → Goal



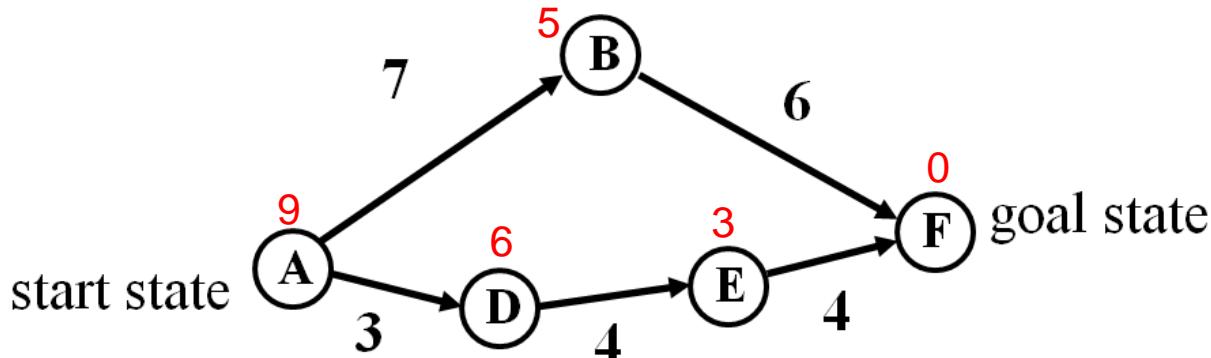
Greedy Best-first Search

Completeness

- Greedy best first search algorithm is complete as in worst case scenario it will search the whole space (worst option).

Optimality

- Greedy best first search algorithm is not optimal.
- A solution can be found in a longer path (lower $h(n)$ with a higher $g(n)$ value).

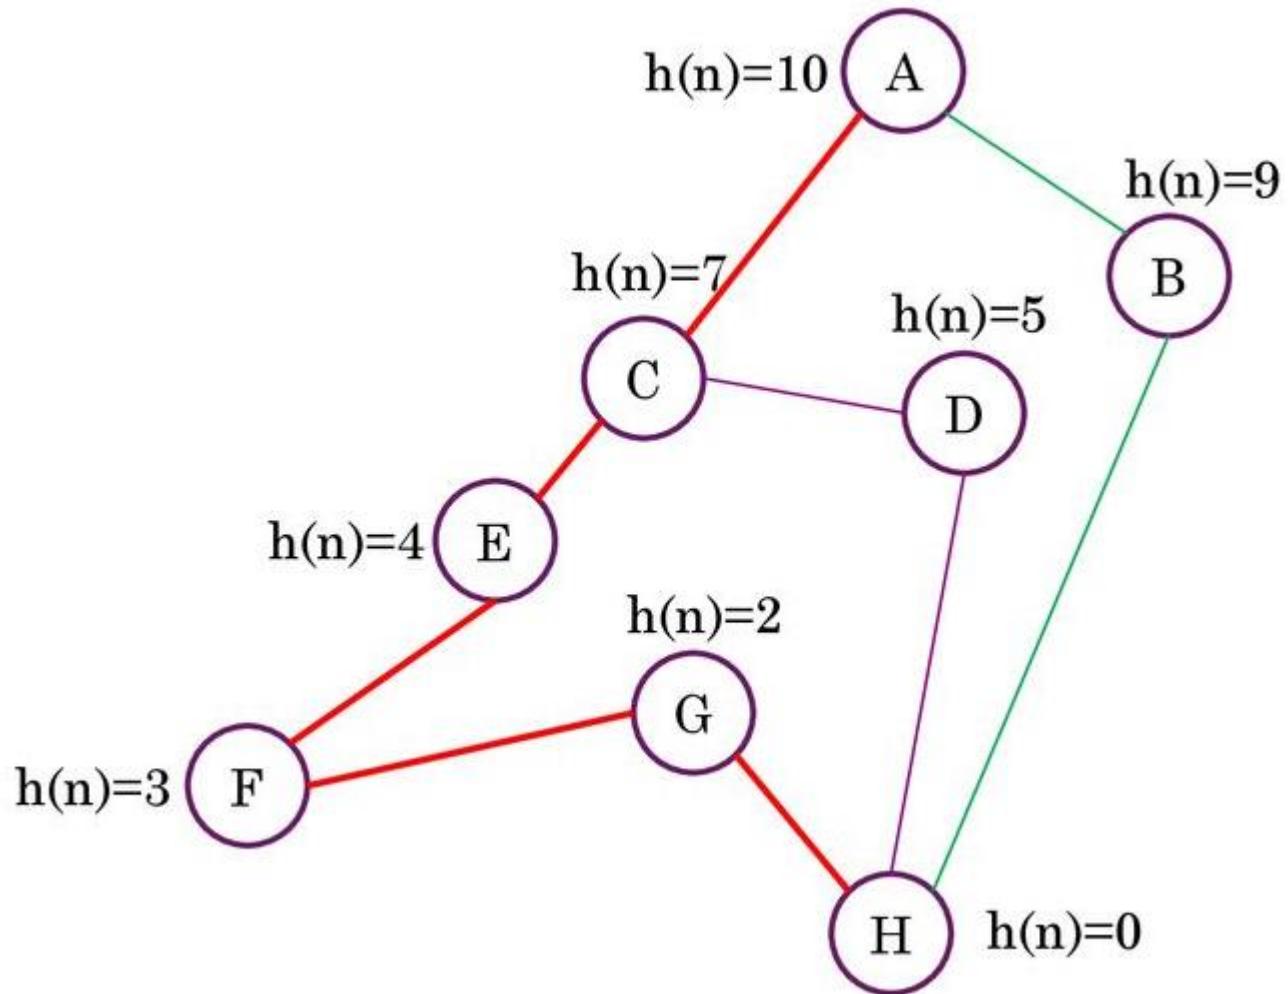


- Say: $h(A) = 9$, $h(B) = 5$, $h(D) = 6$, $h(E) = 3$, $h(F) = 0$
- It does not take the actual cost into account

Greedy Best-first Search

Another example

- The greedy best first search will find the solution path:
 $A \rightarrow C \rightarrow E \rightarrow F \rightarrow G \rightarrow H(\text{Goal})$
- However, this path is the longest path in the graph to the goal and processes more nodes.



A* Search Algorithm

- A* Algorithm is one of the best and popular techniques used for path finding and graph traversals.
- A lot of games and web-based maps use this algorithm for finding the shortest path efficiently.
- It combines features of UCS and greedy best-first search, by which it solve the problem efficiently.
- It uses the evaluation function $f(n) = g(n) + h(n)$
 - $g(n)$ is the cumulative cost from the start node to the node n
 - $h(n)$ is the heuristic (estimated) cost from the node n to the goal
 - $f(n)$ is the estimated total cost of path through n to the goal

A* Search Algorithm

Advantages:

- The A* algorithm can obtain the best solution because it considers both the cost calculated up to now and the estimated future cost.
- A* search algorithm is optimal and complete.
- This algorithm can solve very complex problems.

Disadvantages:

- A* search algorithm has some complexity issues.
- The main drawback of A* is memory requirement as it keeps all generated nodes in the memory, so it is not practical for various large-scale problems.

A* Search Algorithm

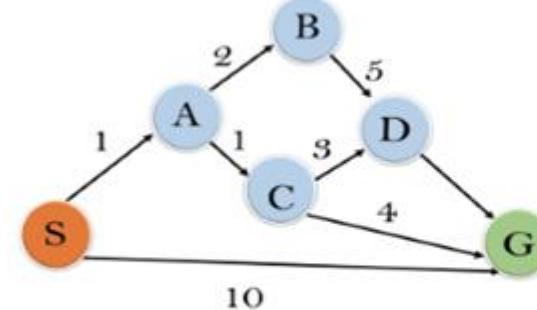
Steps

- Enqueue nodes by their evaluation function $f(n) = g(n) + h(n)$.
- Sort nodes by ascending value of f
- If there is more than one path to the same node, keep only the path with the lowest cost in the queue.
- Stop when the goal is the first node in the queue
Else, expand nodes starting from the first node in the queue
- Ignore back paths to the start node.
- Ignore back paths to the parent node.

A* Search Algorithm

Example

- In this example, we will traverse the given graph using the A* algorithm.
- The heuristic value of all states is given in the table.



State	$h(n)$
S	5
A	5
B	4
C	2
D	6
G	0

A* Search Algorithm

Solution

Expanded nodes list

- S { f(A), f(G) } not sorted

$$f(A)=g(A)+h(A)=1+5=6$$

$$f(G) = g(G)+h(G) = 10+0=10$$

S { A6, G10 } sorted

- A6 { G10, f(B), f(C) } not sorted

$$f(B)=g(B)+h(B)=3+4=7$$

$$f(C) = g(C)+h(C) = 2+2=4$$

A6 { C4, B7, G10 } sorted

- C4 { B7, G10, f(D), f(G)} not sorted

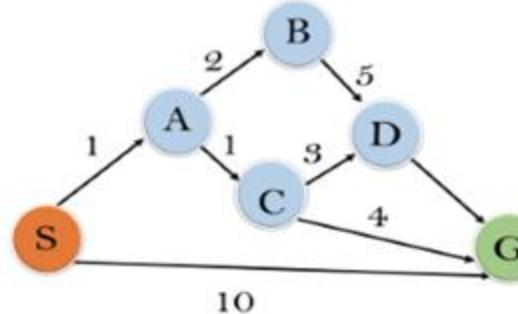
$$f(D)=g(D)+h(D)=5+6=11$$

$$f(G) = g(G)+h(G) = 6+0=6$$

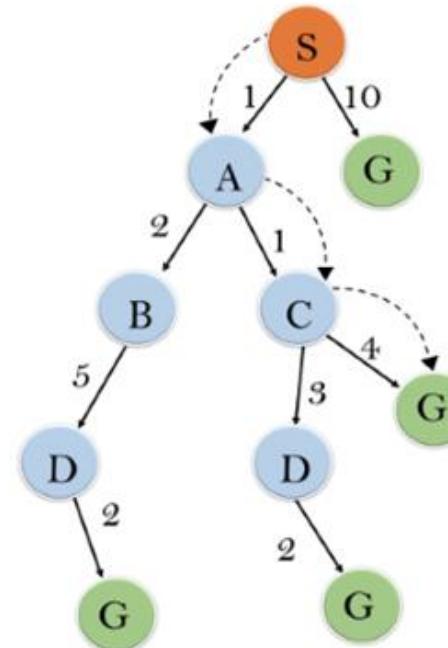
C4 { G6, B7, G10, D11 } sorted

- G6

- Solution path found is S → A →C → G



State	h(n)
S	5
A	5
B	4
C	2
D	6
G	0



A* Search Algorithm

Practical Examples: 8-Puzzle

- Given an initial state of an 8-puzzle problem and final state to be reached.
- Find the most cost-effective path to reach the final state from initial state using A* Algorithm.
- Consider the cost from one state to another = 1 (which will be used to get $g(n)$), and $h(n) = \text{Number of misplaced tiles}$.

2	8	3
1	6	4
7		5

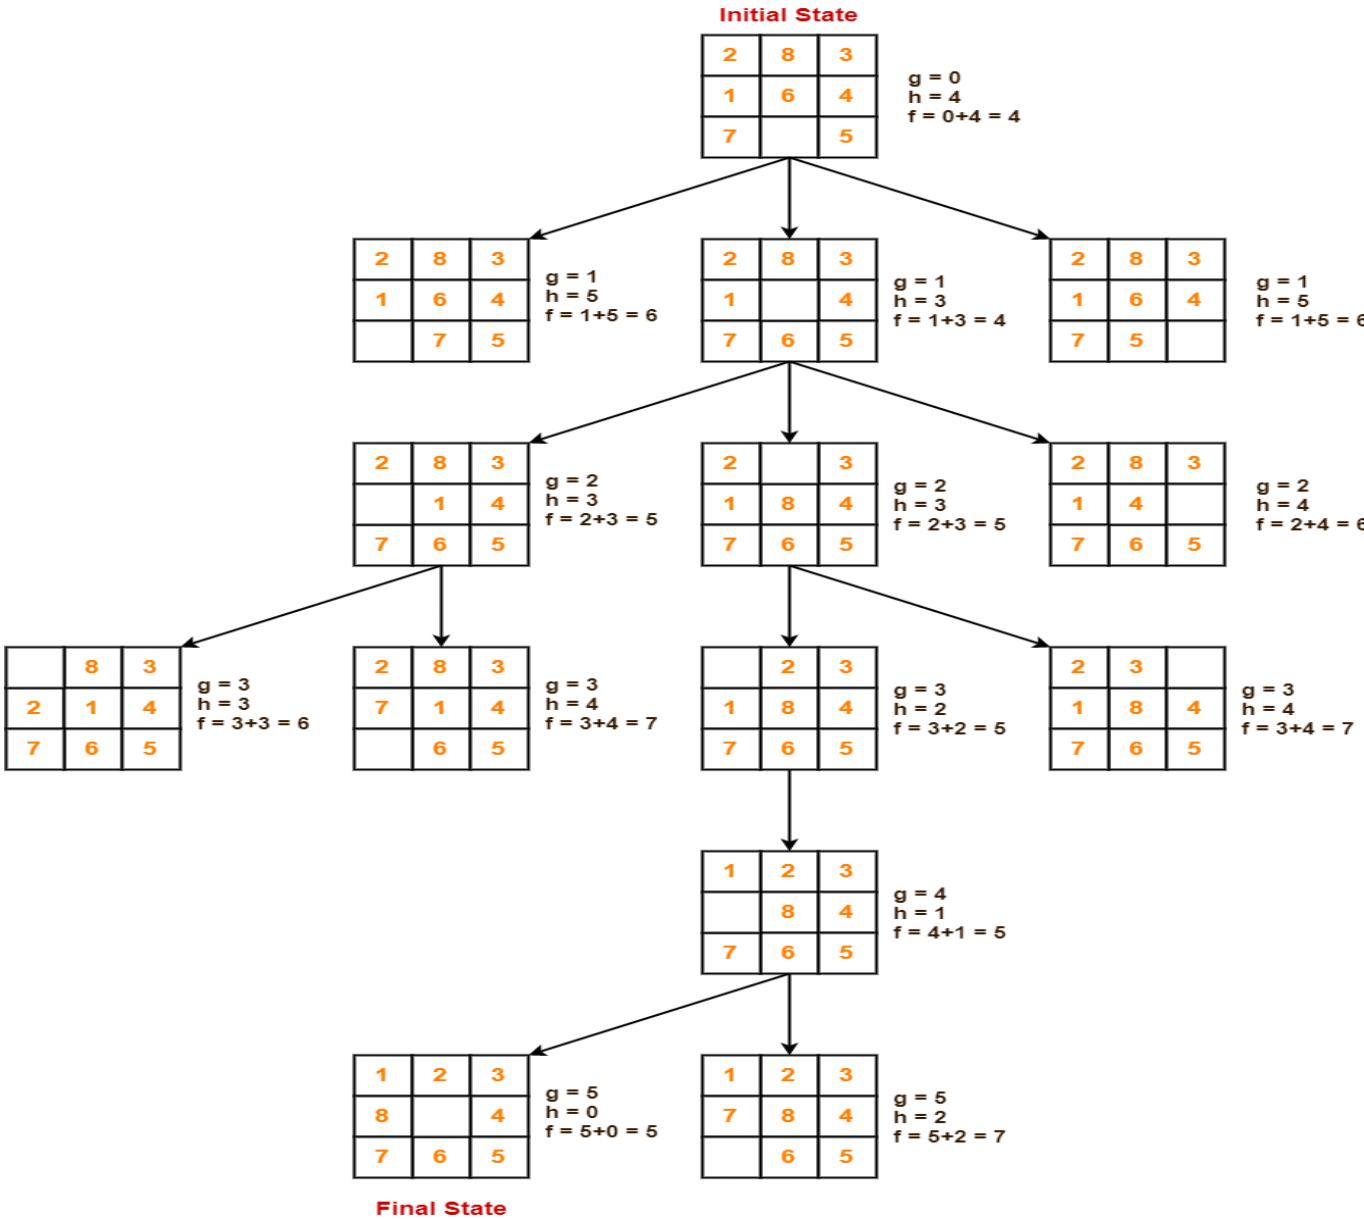
Initial State

1	2	3
8		4
7	6	5

Final State

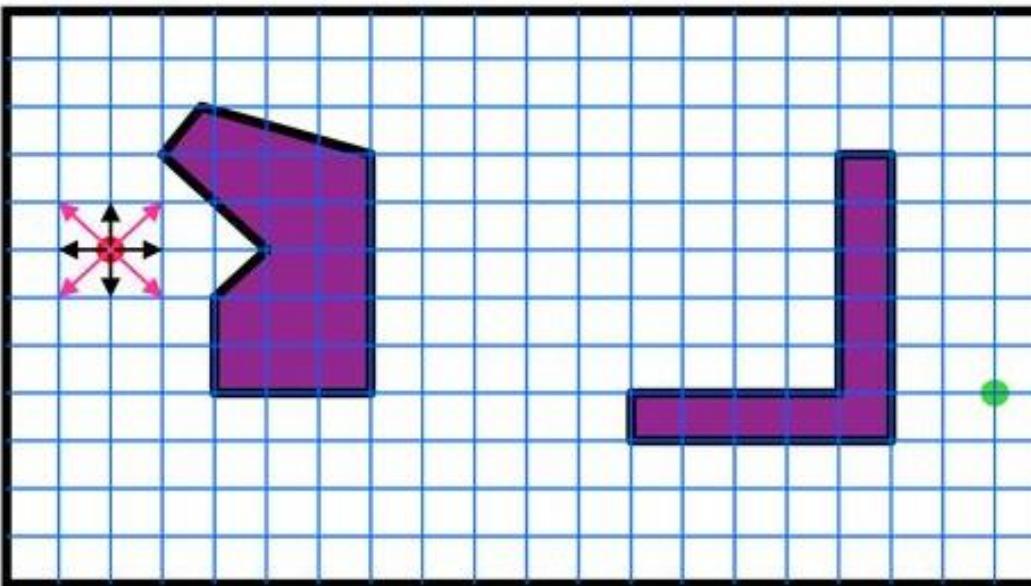
A* Search Algorithm

Solution



A* Search Algorithm

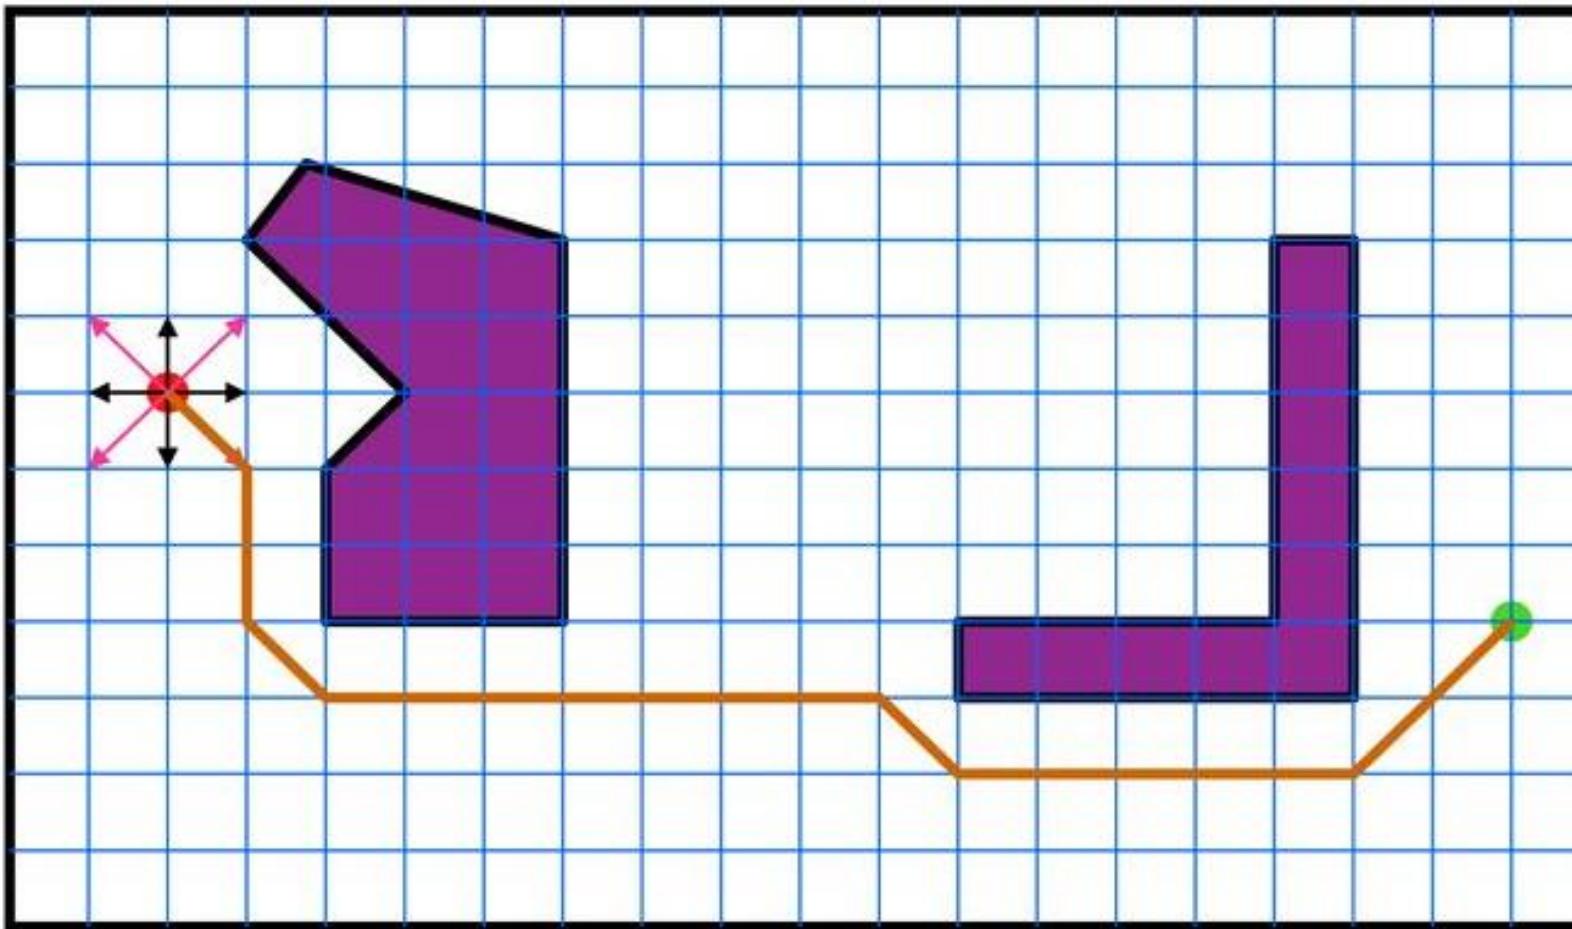
Practical Examples: Robot Navigation



- $f(n) = g(n) + h(n)$, with $h(n)$ = straight-line distance from node n to goal
- $h(n) = \sqrt{(x_g - x_n)^2 + (y_g - y_n)^2}$
- $g(n)$: Cost of one horizontal/vertical step = 1
Cost of one diagonal step = $\sqrt{2}$

A* Search Algorithm

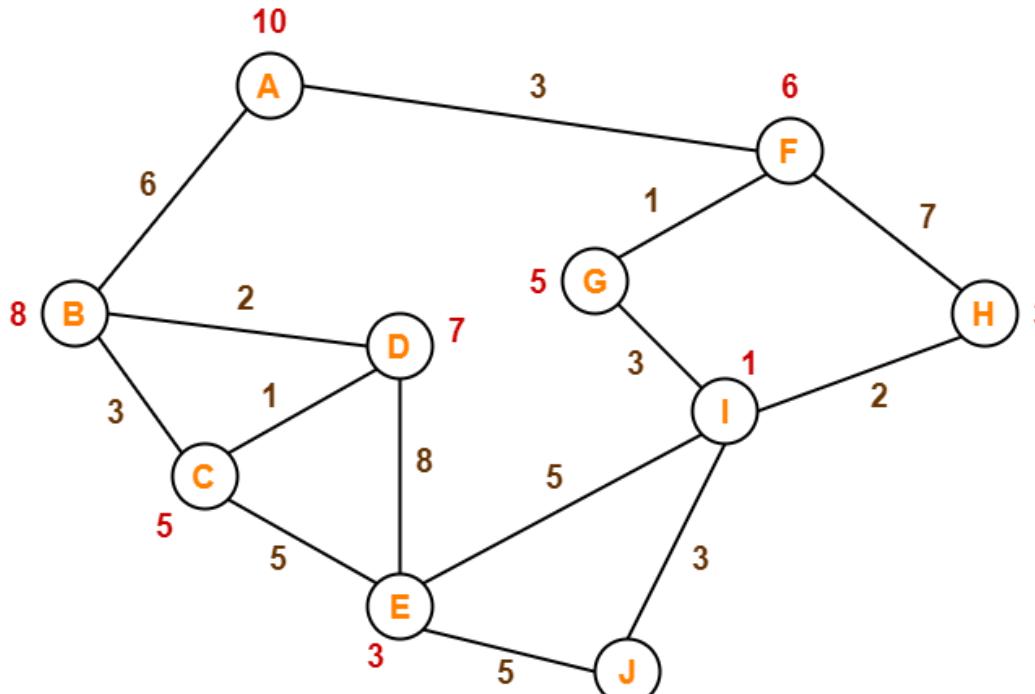
Solution: Robot Navigation



A* Search Algorithm

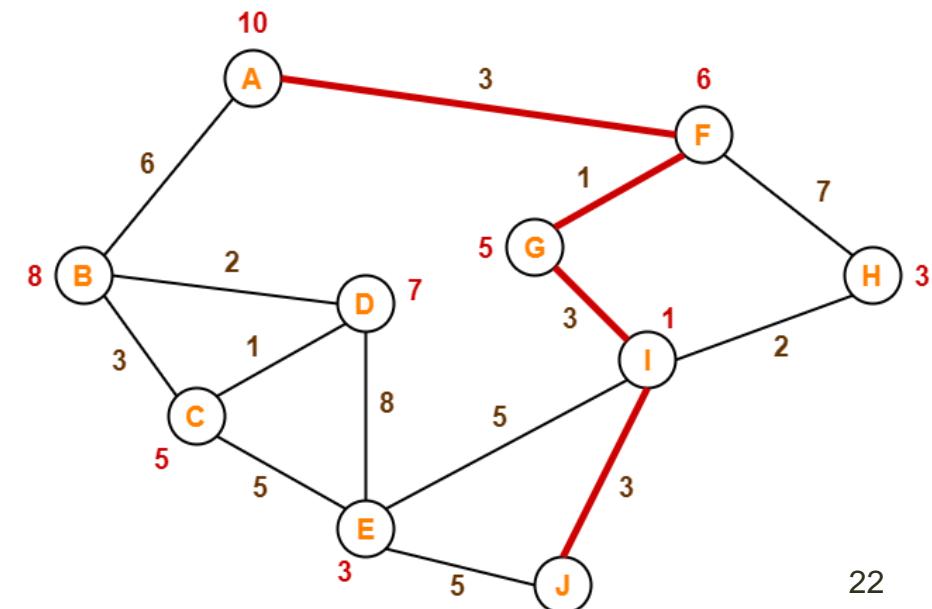
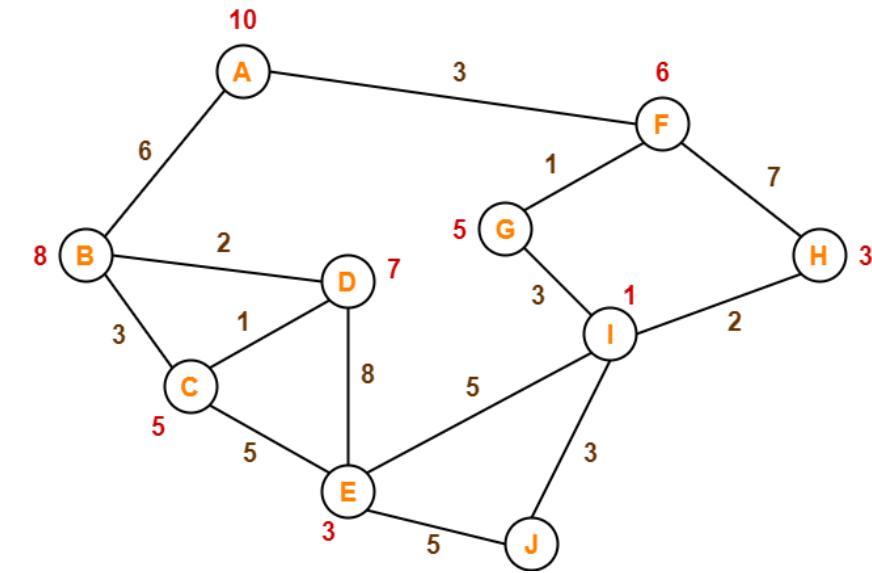
Practical Examples: Map

- Consider the following map of cities
- The numbers written on edges represent the distance between the cities.
- The numbers written on nodes represent the straight-line distances to the target city.
- Find the most cost-effective path to reach from start city A to target city J using A* Algorithm.



A* Search Algorithm

Solution



A* Search Algorithm

Practical Examples: 8-Puzzle

- Given an initial state of an 8-puzzle problem and final state to be reached.
- Find the most cost-effective path to reach the final state from initial state using A* Algorithm. Assume the cost of one move = 1

5		8
4	2	1
7	3	6

N

1	2	3
4	5	6
7	8	

goal

- You can take the heuristic function $h(n)$ as:
 - Number of misplaced tiles
 - Sum of distances of each tile to goal

